

View Maintenance in Data Warehousing

CV Hariharan •1610110147

Divya Raj •1610110123

Devansh Purohit •1610110116

Acknowledgment

We would like to thank Dr. Sonia Khetarpaul, our instructor for the Advance Database Management course this semester, for helping us with this project and the report. She taught us the concepts needed for this project and was available to us for help, whenever we needed it. Thank you.

We would also like to thank Hemant Jain and Anjana Gosain for their amazing paper titled "A Comprehensive Study of View Maintenance Approaches in Data Warehousing Evolution". It was a source of great help throughout this project.

Also, huge gratitude towards Abdulaziz S. Almazyad and Mohammad Khubeb Siddiqui for their paper titled "Incremental View Maintenance: An Algorithmic Approach". It gave us great insight into the incremental approach to view maintenance and inspired the final solution that we implemented.

Finally, we would like to thank Github user AntonioL for his repository - "factorized-incremental-maintenance". It provided us a code wise basis for the project and helped us get started.

All the relevant links can be found at the end of this report in the Bibliography section.

Contents

1	Introduction	3
2	Existing and Related Work	4
3	Proposed Approach	6
4	Implementation and Results	8
5	Limitations and Conclusions	10

1 Introduction

A data warehouse mainly stores integrated information over data from many different remote data sources for query and analysis. The integrated information at the data warehouse is stored in the form of materialized views. Using these materialized views, user queries may be answered quickly and efficiently as the information may be directly available. These materialized views must be maintained in answer to actual relation updates in the different remote sources.

One of the issues related to materialized views is that whether they should be recomputed or they should be adapted incrementally after every change in the base relations.

View maintenance is the process of updating a materialized view in response to changes to the underlying data is called view maintenance. There are several algorithms developed by different authors to ease the problem of view maintenance for data warehouse systems. [2]

In this report, we document the already existing approaches in the field, describe the approach that we took and also show our findings and conclusions.

First, some definitions-

- **SOURCE** - A database, application, file, or other storage facility from which the data in a data warehouse is derived. The source contains the operating data, flat files and stage files. The stage file receives the data from source process and it verifies its credit-ability and the required data files will be passed to warehouse through view manager. Source division also termed as top tier of architecture. [1]
- **WAREHOUSE** - A relational database that is designed for query and analysis rather than transaction processing. A data warehouse usually contains historical data that is derived from transaction data, but it can include data from other sources. It separates analysis workload from transaction workload and enables a business to consolidate data from several sources. It contains the Summary data, raw data, metadata, mined data etc. Warehouse division also termed as middle tier of architecture. [1]
- **USER** - Users may be end users and make use of the data warehouse view maintenance in the Analysis of Data mining, Data reporting etc, and User division also termed as top tier of the architecture. [1]

2 Existing and Related Work

Various approaches have been introduced for maintaining the view in a warehouse environment.




Figure 1: Basic Approach

- **BASIC ALGORITHM** In Fig 1, it is shown that there is communication between Source and the warehouse, when update occurs at source, it sends the notification to warehouse later on warehouse sends the query to source for the corresponding update as source receives the query it sends the answer to warehouse to that corresponding query.

1. When an update occurs at the source, it sends the update notification to the warehouse.
2. Warehouse receives the notification and sends back the query to the source about the update.
3. Source receives the query sent by the warehouse and returns the answer to that query.

The basic algorithm is neither convergent nor weakly consistent in warehouse environment. [1]

- **RECOMPUTE VIEW** RV does not rely on incremental view maintenance approach. It is based on recomputation of materialized view from the scratch. When ever the update occurs at the source it recomputes the view from the scratch. In RV approach warehouse sends the Query to the source asking it to recompute the view from the scratch after certain number of updates. RV sends 2 messages for each update. The bytes transferred are much higher in RV than the relative algorithms. This degrades the performance of RV [1]

- **EAGER COMPENSATING ALGORITHM** COLLECT =

W up_i : receive U_i ;

Let $Q_i = v(U_i) - Q \in_{UQS} Q_j(U_i)$
 send Q_i to the source;
 trigger event S qu_i at the source

W ans_i : receive A_i ;

let COLLECT = COLLECT + A_i ;
 if UQS =
 then MV \leftarrow MV + COLLECT; COLLECT \leftarrow
 else do nothing.

ECA is an incremental view maintenance algorithm. It is a method for fixing the view maintenance problem that occurs due to the decoupling between base data and the view maintenance manager at the warehouse. The key idea of the ECA algorithm is that it cannot rely on the state of the base information that is continuously being updated/modified by the sources. It must keep track of the updates received from the source and then filter out i.e., compensate any information that will duplicate the resulting queries. By subtracting (or adding) the results it knows that will (not) get in future queries, it will create an accurate end result for the view.

The above algorithm states that: Initially the COLLECT will be empty, source executes an update (U_i) and the notification sent to the warehouse. Warehouse receives the source update (U_i) and sends the query (Q_i) based on (U_i), for each query in UQS(Unanswered Query Set: the set of query set that were sent by the warehouse, but answers have not been received) formulates a compensating Query Q_j based on U_i and Q_i with Q_j . Warehouse receive the query result and update the Materialized View(MV), the result of the query should be applied to the Materialized View(MV) only after the answer to this query and all related compensating query have been received. To avoid invalid state ECA collects the intermediate answers in relation denote d as COLLECT (initially its empty).[1]

- **LAZY APPROACH** Lazy approach maintains the view in a lazy manner that relieves the updates of the maintenance overhead as in the incremental view maintenance approaches. View maintenance is postponed until the system has free cycles or it is referenced by any query. These free cycles are utilized for the view maintenance that relieves the updates and queries from the overhead. The updates are combined from different transactions

into a single maintenance task. It also exploits row versioning. In lazy maintenance the updates do not maintain the view it just stores the required information so that the affected views can be maintained later. It actually uses system free cycles to maintain the views, in this no updates or queries pay for the maintenance task. But, in case the view is not up to date and query is sent over it, then the particular query has to pay for all part of the view maintenance and some delay also. However, it pays only the view maintenance that it uses and not for other views [2]

3 Proposed Approach

We made an approach which is a hybrid of the incremental and the recompute view approach.

Consider a base relation $P(a, b, c)$ with a materialized view $V(A, B)$, where $V.A$ maps to $P.a$ and $V.B$ maps to $P.b$.




Figure 2: Mapping between base table and the view

1. When the view V is made, 3 triggers are made on P , for insert, update and delete, which will be used later to update the views.
2. Upon inputting the query, the query is first parsed to get the corresponding views, base tables and the operations to be performed.
3. Meanwhile, a query is made to get the mapping between the source's and the warehouse's column. For example, in the case of P and V , to figure out that $V.A = P.a$
4. If the query on the base table is simple (i.e does not contain joins, only selection and projections) go to 6, else go to 5.

5. Simple query - **Incremental approach**, i.e send a command to the view to recompute itself, i.e delete the previous existence, and recalculate the base table query that forms it. Save it to disk. End.
6. Complex query - **Recompute View**, i.e send a query to the view which updates the view instead of recalculating it. This happens by the insert, update and delete triggers that were defined earlier on the base table. End.

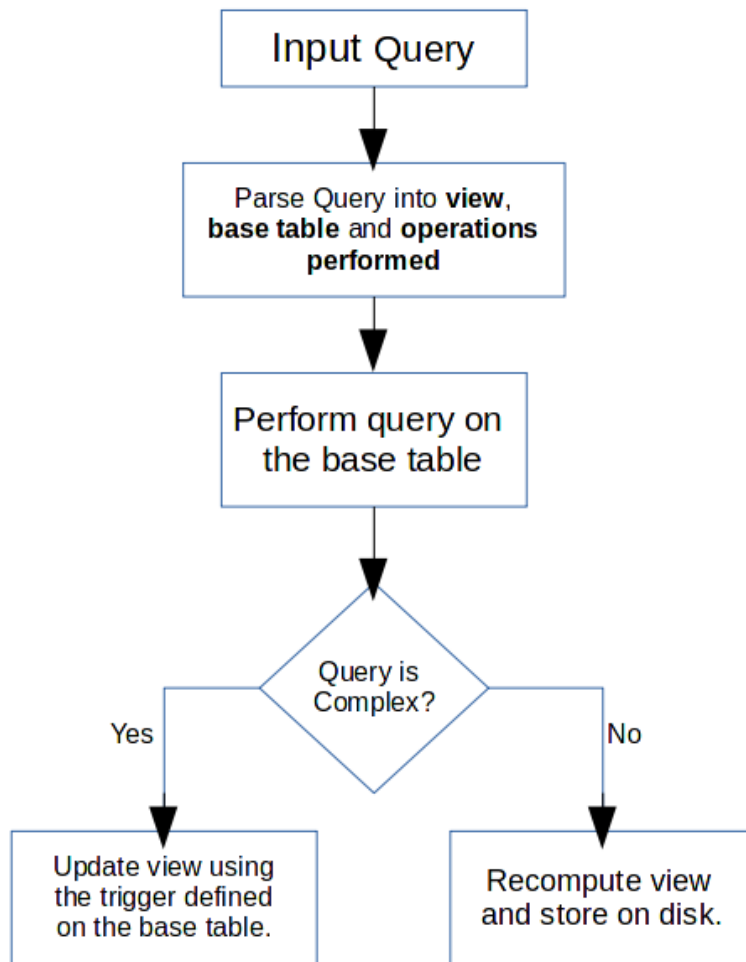


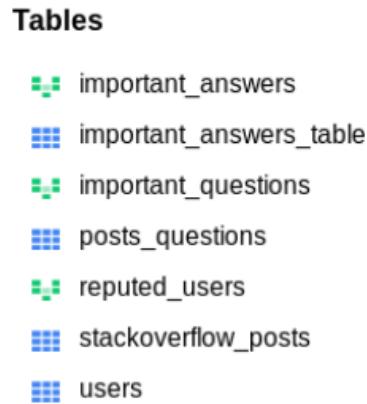
Figure 3: Basic flowchart of the implementation

4 Implementation and Results

To implement our algorithm, we worked with Google Big Query, which is Google’s serverless, highly scalable, enterprise data warehouse designed to make our data analysis on a warehouse without worrying about the logistics of it all. [3]

We wrote our code in the python programming language and used the ‘bigquery’, ‘sqlparse’ and ‘mysql’ libraries to assist us.

We used the openly available StackOverflow dataset as our dummy data warehouse. The dataset is of size 57 GB and has 4 tables and 3 view, as shown in Figure 4.










Tables	
	important_answers
	important_answers_table
	important_questions
	posts_questions
	reputed_users
	stackoverflow_posts
	users

Figure 4: Tables and View in the Dataset (Blue = table, Green = view)

1. Connect to the GCP client using the API key provided, so that we can work with the StackOverflow Dataset.
2. Input a query. Parse it to find out the view, base table(s) and the SQL command involved.
3. If a new view is to be formed from a base table, say ‘important_answers_table’, set up triggers for the base table-
 - (a) Insert- Set up an insert trigger which sends an appropriate query to the view when an insertion of a new answer takes place on the base table.
 - (b) Update- An update trigger which sends an update query to the view when values are updated in the base table.

- (c) Delete- A delete trigger which deletes the corresponding answer from the view as it was deleted in the base table.
4. Figure out the mappings from view to base table. An example result output of step 3 and 4 is shown in Figure 5.

```
<mysql.connector.connection.MySQLConnection object at 0x7faa57e2b0f0>
Query: create table mv as select a, b, c as C from P limit 100
View: mv
Mapping from MV -> Table: {"a": "a", "b": "b", "C": "c"}
INSERT TRIGGER: create trigger trig P ins after insert on P for each row begin insert into mv ( a, b, C) values ( NEW.a, NEW.b, NEW.c);
d;
DELETE TRIGGER: create trigger trig P del after delete on P for each row begin delete from mv where a = OLD.a; end;
UPDATE TRIGGER: create trigger trig P upd after update on P for each row begin update mv set  a = NEW.a, b = NEW.b, C = NEW.c; end;
```

Figure 5: View mapping and trigger formation output

5. Execute the query on the base table, save the result.
6. Based on if the query has more than one tables mentioned after 'from' in the statement. it's classified as simple(1 table) or complex(>1 table). If a query is simple go to step 7 else execute step 8 and then go to step 9.
7. Query is simple, hence the corresponding trigger executes and the value is either inserted, updated or deleted in the view, changes are saved to disk.
8. Query is complex, hence the existing view is deleted and a new view is recalculated. This new view is then saved.
9. End.

5 Limitations and Conclusions

The approach that we've taken is really basic and has a few limitations.

The obvious one is the recomputation of the view the query made is complex i.e contains joins etc. That is one thing we would want to improve in the future. This can be done by having a logging system between the view and the base table, which can be used to calculate the diffs instead of the entire view again. Another limitation is that our solution is not as scalable as some of the existing approaches (notably ECA). It runs smoothly on our StackOverflow dataset, but will mostly struggle in terms of performance when used on bigger and more varied datasets.

In conclusion, we would like to say that even though we were not able to implement some of the advanced view maintenance algorithms, this project gave us an opportunity to study about these algorithms thoroughly, and develop a theoretical understanding about them.

The actual implementation that we did, gave us good insight into how materialized views are maintained practically. This gave us an opportunity to extend what we had learned in the class and build something practical that could reinforce our learning of the concept.

Bibliography

- [1] Abdulaziz S. Almazyad and Mohammad Khubeb Siddiqui *Incremental View Maintenance: An Algorithmic Approach*. Feb, 2016.
- [2] Hemant Jain Anjana Gosain *A Comprehensive Study of View Maintenance Approaches in Data Warehousing Evolution* . Sept, 2012.
- [3] Google Big Query Documentation <https://console.cloud.google.com/big-query/docs>.
- [4] AntonioL github.com/AntonioL/factorized-incremental-maintenance